

1. Architecture Patterns for Information Retrieval & Semantic Search

There are three dominant architectural patterns for building search systems today:

- **Pattern A: Bi-Encoder (Dense Retrieval)**
 - *Mechanism:* Documents and Queries are independently converted into vectors (embeddings) using a model like BERT. Similarity is calculated using Cosine/Dot Product.
 - *Architecture:* Query \rightarrow Embedding Model \rightarrow Vector DB (ANN Search) \rightarrow Results.
 - *Pros:* Extremely fast retrieval (milliseconds).
- **Pattern B: Cross-Encoder (Re-Ranking)**
 - *Mechanism:* The Query and Document are fed into the model *together* as a pair. The model outputs a relevance score (0 to 1).
 - *Architecture:* Initial Retrieval (Bi-Encoder/BM25) \rightarrow Top 50 Docs \rightarrow Cross-Encoder \rightarrow Top 5 Docs.
 - *Pros:* Higher accuracy than Bi-Encoders because the model sees the interaction between query and document terms.
- **Pattern C: Hybrid Search (The Gold Standard)**
 - *Mechanism:* Run **Sparse Retrieval** (BM25 for keywords) AND **Dense Retrieval** (Vectors for meaning) in parallel.
 - *Architecture:* Merge the results using Reciprocal Rank Fusion (RRF) to get the best of both worlds.

2. Why is it important to have very good search?

- **Conversion/Retention:** In e-commerce, if a user can't find a product, they can't buy it. Bad search directly correlates to lost revenue.
- **Knowledge Worker Productivity:** Employees spend ~20% of their time just looking for information. Good search drastically reduces this wasted time.
- **Trust (RAG):** In LLM applications, the "Generation" is only as good as the "Retrieval." If the search retrieves irrelevant documents, the LLM will hallucinate or fail.

3. How can you achieve efficient and accurate search results in large-scale datasets?

- **Hierarchical Indexing:** Use **HNSW** (Hierarchical Navigable Small Worlds) graphs. This allows logarithmic time complexity $O(\log N)$ search instead of linear scan.
- **Two-Stage Retrieval:**
 1. **Stage 1 (Recall):** Retrieve 1,000 candidates using a fast, approximate index (HNSW/IVF).
 2. **Stage 2 (Precision):** Re-rank the top 50 using a heavy Cross-Encoder model.

- **Quantization:** Compress vectors (e.g., from float32 to int8). This reduces memory usage by 4x and speeds up distance calculations, allowing you to fit billions of vectors in RAM.
-

4. Scenario: Improving a Failing RAG Retrieval System

Problem: The RAG system is retrieving irrelevant chunks.

Steps to Improve:

1. **Analyze Failures:** Look at the bad queries. Is it missing keywords (Product ID "X-500") or missing semantic meaning?
 2. **Hybrid Search:** If it misses keywords, add **BM25** (keyword search) alongside the vector search.
 3. **Chunking Strategy:** Is the chunk size too small (missing context) or too big (too much noise)? Experiment with **Parent-Child Retrieval** (search small chunks, retrieve the parent document).
 4. **Query Expansion:** Use an LLM to rewrite the user's query.
 - *User: "It's slow." \$\rightarrow\$ Rewrite: "Why is the dashboard loading latency high?"*
 5. **Re-Ranking:** Implement a Cross-Encoder (like bge-reranker) to filter the top 50 retrieved results before sending them to the LLM.
-

5. Explain the keyword-based retrieval method.

- **Mechanism:** It matches exact words in the query to words in the documents. The industry standard is **BM25** (Best Matching 25).
 - **How BM25 works:** It isn't just counting words (TF). It penalizes common words (IDF) and normalizes for document length.
 - *TF (Term Frequency):* How often does "apple" appear in this doc?
 - *IDF (Inverse Document Frequency):* Is "apple" rare across the whole database? (If yes, it's more important).
 - **Pros:** Exact matching (great for names, IDs, error codes). Zero hallucinations.
 - **Cons:** Fails at synonyms. "Car" will not match "Automobile."
-

6. How to fine-tune re-ranking models?

You don't fine-tune the vector database; you fine-tune the **Cross-Encoder**.

1. **Dataset:** Create pairs of (Query, Document, Label).
 - *Label:* 1 (Relevant) or 0 (Irrelevant).
2. **Hard Negatives:** Mine "Hard Negatives"—documents that the base model *thought* were relevant (high score) but are actually wrong.
3. **Training:** Use a **Contrastive Loss** or **Cross-Entropy Loss** function. The model learns to output a score closer to 1 for the positive pair and 0 for the negative pair.
4. **Library:** Use sentence-transformers **CrossEncoder** class for easy implementation.

7. Most common metric in Information Retrieval and when it fails?

- **Metric: Recall@K** (Did the correct answer appear in the top K results?).
- **Failure Case:** Recall only cares *if* the answer is there, not *where* it is.
 - *Scenario:* If the correct answer is at Rank #10 (bottom of page 1), Recall@10 is 100%. But the user likely clicked the first result and left. Recall fails to measure the **ranking quality**.

8. Metric for Quora-like System (Pertinent Answers Quickly)

- **Choice: MRR (Mean Reciprocal Rank).**
- **Why:** You want the *best* answer to be at the very top (Rank 1).
 - If the answer is at Rank 1, Score = 1.
 - If at Rank 2, Score = 1/2 (0.5).
 - If at Rank 10, Score = 1/10 (0.1).
- MRR heavily penalizes the system if the "pertinent" answer drops even slightly down the list. This aligns with the goal of "finding answers quickly."

9. Metric for a Recommendation System?

- **Choice: NDCG (Normalized Discounted Cumulative Gain).**
- **Why:** Unlike search (where there is usually *one* right answer), recommendations have *multiple* good items with varying degrees of relevance (Perfect match, Good match, Okay match).
- **How it works:** NDCG gives credit for retrieving relevant items but "discounts" the credit if they are lower down the list. It handles graded relevance better than MRR or Recall.

10. Comparing IR Metrics

Metric	Focus	Use When...
Recall@K	"Is the answer in the list?"	You just need to find the document <i>somewhere</i> (e.g., Legal Discovery).
Precision@K	"Is the list mostly junk?"	You want to avoid showing bad results (e.g., Google Search first page).

MRR	"Is the answer at the top?"	There is one correct answer (Factoid QA, Quora).
NDCG	"Is the ranking order perfect?"	There are multiple relevant items with different quality levels (Netflix, E-commerce).

11. How does Hybrid Search work?

Hybrid search combines the strengths of **Keyword Search (BM25)** and **Semantic Search (Vectors)**.

1. **Run Sparse Search (BM25):** Finds documents with exact keywords (e.g., "Error 404").
 2. **Run Dense Search (Vectors):** Finds documents with similar meaning (e.g., "Page not found").
 3. **Normalization:** The scores from BM25 (e.g., 0 to 15) and Vectors (e.g., 0.6 to 0.9) are on different scales. You must normalize them (usually 0 to 1).
 4. **Fusion:** Combine the two lists using a weighted sum $(0.7 * \text{Vector}) + (0.3 * \text{BM25})$ or RRF.
-

12. How would you merge and homogenize rankings from multiple methods?

- **Technique: Reciprocal Rank Fusion (RRF).**
- **Why:** It doesn't rely on the raw scores (which might be incompatible between models). It relies only on the **Rank**.
- **Formula:** For each document, calculate a score:

$$\text{Score} = \sum \frac{1}{k+1} \text{rank}$$
- **Logic:** If Doc A is Rank 1 in Keyword Search and Rank 1 in Vector Search, it gets a massive score. If Doc B is Rank 1 in Vector but Rank 100 in Keyword, its score drops significantly. This effectively "homogenizes" the lists without needing to calibrate the model scores.

13. How to handle multi-hop/multifaceted queries?

- **Problem:** Query: "Who was the CEO of the company that created the iPhone?"
 1. **Hop 1:** Who created the iPhone? \rightarrow Apple.
 2. **Hop 2:** Who was the CEO of Apple? \rightarrow Steve Jobs.
- **Technique: Query Decomposition.**
 1. Use an LLM to break the complex query into sub-questions.
 2. **Step 1:** Retrieve docs for "Who created iPhone?". Result: "Apple".

3. **Step 2:** Use the answer from Step 1 to generate a new query: "Who was the CEO of Apple?".
 4. **Final:** Retrieve docs for Step 2 and synthesize the answer.
-

14. What are different techniques to be used to improved retrieval?

1. **Metadata Filtering:** Filter by Year > 2023 before searching vectors.
2. **HyDE (Hypothetical Document Embeddings):** Generate a fake answer, embed *that*, and search.
3. **Query Expansion:** Add synonyms to the query (e.g., "Car" \rightarrow "Car, Auto, Vehicle").
4. **Contextual Re-Ranking:** Use a Cross-Encoder to re-score the top results.
5. **Sentence Window Retrieval:** Retrieve a single sentence for the search match, but give the LLM the 5 sentences before and after it for context.