

1. What is a vector database?

A **Vector Database** is a specialized database designed to store, manage, and index **high-dimensional vector embeddings**.¹ Unlike a traditional database that stores rows and columns of text or numbers, a vector database stores data as mathematical vectors (lists of floating-point numbers) derived from embedding models.²

+1

Its primary function is to perform **Vector Similarity Search**—finding items that are "closest"³ to a query vector in 3D or N-dimensional space.

2. How does a vector database differ from traditional databases?

Feature	Traditional Database (SQL/NoSQL)	Vector Database
Data Structure	Rows, Columns, Key-Value pairs, JSON.	High-dimensional Vectors (Embeddings).
Search Method	Exact Match: WHERE id = 101 or WHERE text LIKE '%apple%'. 	Similarity Search: "Find data semantically similar to 'apple'".
Logic	Boolean (True/False). A record either matches or it doesn't.	Probabilistic/Distance-based. Returns a ranked list of "nearest neighbors."
Use Case	Transactional apps, inventory, user profiles.	RAG, Recommendation Systems, Semantic Search, Anomaly Detection.

3. How does a vector database work?

The workflow consists of three main stages:

1. **Vectorization:** Raw data (text, images) is converted into vectors using an embedding model (e.g., OpenAI, HuggingFace).

2. **Indexing:** The database maps these vectors into a data structure (Index) that allows for fast searching (e.g., HNSW, IVF).⁴ It doesn't just pile them up; it organizes them geometrically.⁵
3. +1
4. **Querying (ANN):** When you search, the DB calculates the distance (Cosine/Euclidean) between your query vector and the stored vectors to find the **Approximate Nearest Neighbors (ANN)**.

4. Vector Index vs. Vector DB vs. Vector Plugins

It is crucial to distinguish between the algorithm and the infrastructure.

- **Vector Index:** The algorithm or data structure used to organize vectors for search (e.g., FAISS, HNSW).⁶ It is just a library, not a full system. It runs in memory and doesn't handle CRUD (Create, Read, Update, Delete) well.
- **Vector Database:** A full management system built around an index.⁷ It handles storage, CRUD operations, filtering, scalability, sharding, and persistence (e.g., Pinecone, Milvus, Qdrant).
- +1
- **Vector Plugin:** An add-on for a traditional database that enables vector search capabilities (e.g., pgvector for PostgreSQL, Elasticsearch with vector search).⁹

5. Scenario: Small Dataset Search Strategy

Scenario: You have a small dataset of customer reviews. You need **perfect accuracy**, and speed is **not** a primary concern.

Strategy: Flat Indexing (Brute Force / Exact Search).¹⁰

- **Why:** A "Flat" index calculates the distance between the query vector and **every single vector** in the database.
- **Pros:** It guarantees 100% Recall (Accuracy). You will never miss the true nearest neighbor.
- **Cons:** It is slow ($O(N)$ complexity). However, since the dataset is **small**, the latency penalty is negligible. Approximate methods (like HNSW) trade off accuracy for speed, which is unnecessary here.

6. Vector Search Strategies: Clustering & LSH

Clustering (IVF - Inverted File Index)

This method partitions the vector space into distinct regions (clusters).

- **How it works:**
 1. Select k center points (centroids) in the vector space.
 2. Assign every vector to its nearest centroid. This creates "cells" (Voronoi regions).
 3. **Search:** When a query comes in, compare it to the centroids first. Identify the closest centroid, and then search *only* the vectors inside that specific cell.¹¹

Locality-Sensitive Hashing (LSH)

A method that hashes similar input items into the same "buckets" with high probability.¹²

- **How it works:** It uses random projections (hash functions) to slice the high-dimensional space.¹³ Vectors that are close to each other are likely to land in the same hash bucket.¹⁴
- +1
- **Search:** You simply check the bucket your query hashes into. It is very fast but often yields lower accuracy (recall) compared to graph-based methods.

7. How does Clustering reduce search space? (Failures & Mitigation)

- **Reduction:** Instead of comparing the query to 1 million vectors, you compare it to 1,000 centroids. If the closest centroid contains 1,000 vectors, you only search those 1,000 + the centroids. This reduces the search from N to \sqrt{N} (roughly).
- **Failure (The Edge Problem):** If a query vector lands near the *edge* of a cluster, its true nearest neighbor might actually be just across the border in an adjacent cluster. Since IVF normally ignores neighboring clusters, it will miss this match.
- **Mitigation:** Use the `nprobe` parameter.¹⁵ Instead of searching only the *single* closest cluster, you search the **top 3 or 5 closest clusters**. This drastically improves accuracy at a slight cost to speed.

8. Random Projection Index

This is a dimensionality reduction technique based on the **Johnson-Lindenstrauss lemma**.

- **Concept:** It projects high-dimensional vectors (e.g., 1536 dims) into a lower-dimensional space (e.g., 64 dims) using a random matrix.¹⁶
- **Result:** It preserves the relative distances between points surprisingly well.
- **Use Case:** Good for initial coarse filtering to speed up calculations before doing a fine-grained search.

9. Product Quantization (PQ) Indexing

PQ is a **lossy compression** technique for vectors.¹⁷

- **How it works:**
 1. **Split:** Break a long vector (e.g., 128 dims) into smaller sub-vectors (e.g., 8 chunks of 16 dims).¹⁸
 2. **Quantize:** Run clustering (k-means) on each sub-vector independently to find "centroids" for that chunk.¹⁹
 3. **Replace:** Replace the actual sub-vector values with the *ID* of the closest centroid.
- **Result:** A vector that took 4KB might now take 64 bytes.
- **Trade-off:** Massive memory savings (95%+) and faster search, but the distance calculations are approximations, reducing accuracy.

10. Comparing Vector Indexes (Scenario Guide)

Index Type	Recall (Accuracy)	Speed	Memory Usage	Best For...
Flat (Brute Force)	100% (Perfect)	Slow	High (Raw Vectors)	Small datasets (<100k); strict accuracy needs.
IVF (Clustering)	Medium-High	Fast	Medium	Medium/Large datasets; balanced needs.
HNSW (Graph)	High	Very Fast	High (Graph overhead)	Real-time search; performance critical apps.

PQ (Quantization)	Low-Medium	Fast	Very Low (Compressed)	Massive datasets (1B+); RAM constrained environments.
------------------------------	------------	------	--------------------------	--

Scenario Decision:

- *Project:* Real-time recommendation engine for 50M items.
- *Choice:* **HNSW**. It offers the best latency/recall trade-off, provided you have the RAM to support the graph structure.

11. How to decide Ideal Similarity Metrics?

1. **Euclidean Distance (L2):** Measures the straight-line distance between points.²⁰
 - *Use when:* Magnitude matters (e.g., Anomaly detection where "location" in space implies "normalcy").
2. **Cosine Similarity:** Measures the **angle** between vectors.²¹ Normalizes magnitude.
 - *Use when:* Text similarity (e.g., Document search).²² A long document and a short document about "Cats" should be considered similar despite different vector lengths.
3. **Dot Product:** Measures magnitude and projection.²³
 - *Use when:* Recommendation systems (e.g., Matrix Factorization) where higher magnitude (rating) + alignment (preference) implies a better match.²⁴

12. Filtering in Vector DB (Pre vs. Post)

Filtering allows you to combine metadata (e.g., category="shoes") with vector search.²⁵

- **Post-Filtering (Naive):** Perform vector search first to get top 100 results \$\\rightarrow\$ Filter out non-shoes.
 - *Challenge:* If your top 100 results are all "shirts", you end up with **zero results** after filtering.
- **Pre-Filtering (Efficient):** Filter the dataset for "shoes" first²⁶ \$\\rightarrow\$ Perform vector search only on "shoes".²⁷
 - *Challenge:* If the filtered subset is huge, brute forcing it is slow.²⁸ If it's small, building an index for just that subset is inefficient. Modern DBs use **filtered HNSW graphs** or bitmaps to handle this dynamically.

13. How to decide the best Vector Database?

- **Self-Hosted vs. Managed:** Do you have a DevOps team? If no, use **Pinecone** or **Weaviate Cloud**. If yes, consider **Milvus** or **Qdrant** (Docker containers).
- **Ecosystem Integration:** If you already use Postgres, **pgvector** is the easiest choice (Tea Kettle Principle).²⁹ If you are deep in the AWS ecosystem, **OpenSearch** is natural.
- **Latency vs. Scale:**
 - Need <10ms latency? **Qdrant** (Rust-based) or **Pinecone**.³⁰
 - Have 1 Billion+ vectors? **Milvus** (Highly scalable, distributed).³¹
- **Features:** Do you need hybrid search (Keyword + Vector)? Look for DBs with strong BM25 integration (Weaviate, Qdrant).