
1. What is the difference between Predictive/Discriminative AI and Generative AI?

- **Answer:**
 - **Discriminative AI:** Models the decision boundary between classes. It learns the conditional probability $P(Y|X)$ (Given input X , what is the probability it belongs to class Y ?).
 - *Examples:* Logistic Regression, SVM, CNNs for classification (Cat vs. Dog).
 - **Generative AI:** Models the distribution of the data itself. It learns the joint probability $P(X, Y)$ or just $P(X)$. It tries to understand *how* the data is created so it can generate *new* samples that look like the training data.
 - *Examples:* GANs, VAEs, and LLMs (Next Token Prediction).
- **Follow-up: Which requires more data to train and why?**
 - **A:** Generative models generally require significantly more data. Learning the entire distribution of data points (to generate new ones) is a much harder mathematical problem than simply finding a line/boundary that separates two groups.

2. What is LLM, and how are LLMs trained?

- **Answer:**
 - An LLM is a deep learning model (usually Transformer-based) trained on massive text corpora to predict the next token in a sequence.
 - **Training Pipeline:**
 1. **Pre-training:** Self-supervised learning on massive datasets (web, books). Objective: Next Token Prediction. Result: Base Model (learns grammar, facts, reasoning).
 2. **Supervised Fine-Tuning (SFT):** Training on high-quality (Instruction, Output) pairs. Result: Instruct Model (learns to follow commands).
 3. **RLHF/DPO (Alignment):** Reinforcement Learning from Human Feedback or Direct Preference Optimization. Result: Chat Model (aligned with human values like safety and helpfulness).
- **Follow-up: What is the difference between Pre-training and Fine-tuning regarding computational cost?**
 - **A:** Pre-training is astronomically expensive (months on thousands of GPUs). Fine-tuning is relatively cheap (hours/days on a few GPUs).

3. What is a token in the language model?

- **Answer:**
 - A token is the fundamental unit of text an LLM processes. It is **not** always a word. It can be a character, a sub-word, or a whole word.
 - Modern LLMs use **BPE (Byte Pair Encoding)**.

- *Rule of Thumb*: 1,000 tokens \approx 750 words.
 - *Example*: "apple" might be 1 token, but "friendship" might be split into "friend" + "ship" (2 tokens).
 - **Follow-up: Why do LLMs struggle with arithmetic/math?**
 - **A:** Tokenization often splits numbers inconsistently. "100" might be one token, but "101" might be "10" and "1". This makes it hard for the model to learn digit-by-digit operations.
-

4. How to estimate the cost of running SaaS-based and Open Source LLM models?

- **Answer:**
 - **SaaS (OpenAI/Anthropic)**: Cost = (Input Tokens \times Price) + (Output Tokens \times Price). *Note: Output tokens are usually 3x-10x more expensive.*
 - **Open Source (Self-Hosted)**:
 - **GPU Cost**: Hourly rate of the GPU (e.g., A100/H100).
 - **VRAM Requirement**: ~2 bytes per parameter (FP16). A 7B model needs ~14GB just to load, plus extra for the KV Cache (context).
 - **Utilization**: Cost depends heavily on batch size. Low utilization = high cost per token.
-

5. Explain the Temperature parameter and how to set it.

- **Answer:**
 - Temperature scales the logits (raw scores) before the Softmax function during decoding.
 - **Low Temperature ($<1.0$$)**: Sharpens the probability distribution. The model picks the most likely token. Result: Deterministic, factual, focused. (Use for: Coding, Math).
 - **High Temperature ($>1.0$$)**: Flattens the distribution. Less likely tokens get a chance. Result: Creative, random, diverse. (Use for: Creative writing, Brainstorming).
-

6. What are different decoding strategies for picking output tokens?

- **Answer:**
 - **Greedy**: Always pick the highest probability token. (Fast, but repetitive).
 - **Beam Search**: Track top K sequences (beams) at each step. (Better quality, computationally expensive).
 - **Top-K Sampling**: Sample from the top K most likely tokens only.

- **Top-P (Nucleus) Sampling:** Sample from the smallest set of tokens whose cumulative probability adds up to \$P\$ (e.g., 0.9). This is **dynamic**—if the model is unsure, the pool is big; if sure, the pool is small.
- **Follow-up:** *Why is Top-P generally preferred over Top-K?*
 - **A:** Top-P adapts to the context. Top-K is rigid (always 50 tokens), which cuts off good options in "flat" distributions or includes bad options in "sharp" distributions.

7. What are different ways you can define stopping criteria in large language model?

- **Answer:**
 1. **Max Tokens:** Hard limit (e.g., stop after 512 tokens).
 2. **Stop Sequences:** Specific strings (e.g., "\n", "User:") that trigger a halt.
 3. **EOS Token:** The model predicts its own <End_Of_Sequence> token.
 4. **Time Limit:** Stop after \$X\$ seconds (latency constraints).

8. How to use stop sequences in LLMs?

- **Answer:**
 - Stop sequences prevent the model from rambling or hallucinating the next part of a conversation.
 - **Example:** In a chatbot, if the prompt format is User: Hi \n AI: Hello, you must set User: as a stop sequence. Otherwise, the AI might generate User: How are you? immediately after its own answer, effectively talking to itself.

9. Explain the basic structure prompt engineering.

- **Answer:**

A robust prompt has four parts:

 1. **Persona/Role:** "Act as a Senior Python Developer."
 2. **Instruction:** "Write a script to scrape a website."
 3. **Context/Constraint:** "Use the BeautifulSoup library. Handle 404 errors."
 4. **Format:** "Output the code in Markdown."

10. Explain in-context learning.

- **Answer:**
 - The ability of an LLM to learn a new task at inference time by seeing examples in the prompt, **without any weight updates**.

- If you show it English: Dog -> French: Chien, it learns the pattern "Translation" and applies it to the next input.
 - **Follow-up:** Does In-Context Learning update the model's weights?
 - **A:** No. It relies on the model's existing internal representations and attention mechanism to copy the pattern.
-

11. Explain types of prompt engineering.

- **Answer:**
 - **Zero-Shot:** No examples ("Translate this").
 - **Few-Shot:** Providing examples (\$k\$-shot) before the query.
 - **Chain of Thought (CoT):** Asking the model to "think step by step".
 - **RAG:** Injecting retrieved documents into the context.
-

12. What are some of the aspects to keep in mind while using few-shot prompting?

- **Answer:**
 1. **Class Balance:** If you give 3 positive examples and 1 negative, the model will be biased toward "Positive" (Majority Label Bias).
 2. **Recency Bias:** LLMs tend to repeat the label of the *last* example provided.
 3. **Diversity:** Examples should cover edge cases, not just simple ones.
-

13. What are certain strategies to write good prompts?

- **Answer:**
 1. **Delimiters:** Use "'''", ---, or XML tags <data> to separate instructions from content.
 2. **Structured Output:** Explicitly ask for JSON or HTML to make parsing easier.
 3. **Positive Constraints:** Tell the model *what to do* ("Use short sentences") rather than *what not to do* ("Don't use long sentences").
-

14. What is hallucination, and how can it be controlled using prompt engineering?

- **Answer:**
 - **Hallucination:** Confidently stating false information.
 - **Controls:**
 1. **Grounding:** "Answer ONLY using the text below. If the answer is not there, say 'I don't know'."
 2. **Chain of Thought:** Asking for reasoning steps reduces logic errors.

3. **Citations:** Asking the model to quote the source text.

15. How to improve the reasoning ability of LLM through prompt engineering?

- **Answer:**
 - **Chain of Thought (CoT):** Force the model to generate intermediate steps.
 - **Zero-Shot CoT:** Just adding "Let's think step by step" improves math/logic performance significantly.
 - **Few-Shot CoT:** Providing examples that include the reasoning trace (e.g., Q: ... A: First, we calculate X, then Y...).

16. How to improve LLM reasoning if your COT prompt fails?

- **Answer:**
 1. **Self-Consistency:** Generate 5 different CoT paths and pick the most common answer (Majority Voting).
 2. **Reflexion:** Ask the model: "Review your previous answer. Did you miss anything? Correct it."
 3. **Least-to-Most:** Break the problem into sub-questions. Ask the model to solve Q1, then use that to solve Q2.
 4. **Tree of Thoughts (ToT):** Explore multiple reasoning branches and "backtrack" if a branch looks unpromising (like a search algorithm).