

1. How to increase accuracy, reliability & make answers verifiable in LLM?

To move an LLM from a "creative writer" to a "reliable expert," you need to constrain its generation and force it to cite evidence.

- **Grounding (RAG):** Instead of relying on the model's internal memory (which hallucinates), provide a "Source of Truth" (documents, database) in the prompt and force the model to answer *only* using that context.
- **Citations:** Instruct the model to cite the specific document or chunk ID for every claim.
 - *Prompt Example:* "For every sentence you generate, append the source ID like [Source 1]."
- **Chain of Thought (CoT):** accuracy improves when the model is forced to explain its reasoning step-by-step before giving the final answer. This reduces logic errors.
- **Self-Consistency / Voting:** Generate 5 answers for the same prompt (with high temperature) and pick the most common answer. This is statistically more reliable than a single generation.
- **Verifiability:** Use "Evan" (Evaluation) frameworks like **Ragas** or **DeepEval**. You can treat the LLM's output as a claim and use a second "Judge LLM" to verify if the claim is supported by the retrieved context.

2. How does RAG work?

Retrieval-Augmented Generation (RAG) is a technique that connects a generative model (like GPT-4) to an external database. It works in three phases:

1. **Retrieval:** When a user asks a question, the system searches a vector database for relevant information (chunks of text) that matches the user's query intent.
2. **Augmentation:** The retrieved chunks are pasted into the prompt as "Context."
3. **Generation:** The LLM receives the prompt containing both the User Query and the Retrieved Context and generates an answer based *only* on that context.

3. What are some benefits of using the RAG system?

- **Up-to-Date Knowledge:** LLMs have a "training cutoff" (e.g., knowledge stops at 2023). RAG allows the model to answer questions about data created *today* (e.g., stock prices, new emails) without retraining.
- **Reduced Hallucination:** By forcing the model to rely on retrieved context, you significantly lower the chance of it making up facts.
- **Source Citation:** Since the system knows exactly which documents were retrieved, it can show the user the original source (e.g., "See Page 5 of the PDF").
- **Data Security:** You can keep your proprietary data in a secure database and strictly control access (ACLs) during the retrieval step, ensuring the LLM respects user permissions.

4. When should I use Fine-tuning instead of RAG?

This is the classic "Knowledge vs. Behavior" trade-off.

Feature	RAG (Retrieval)	Fine-Tuning (Training)
Goal	New Knowledge: The model needs to know facts it wasn't trained on (e.g., your company's Q3 report).	New Behavior/Style: The model needs to talk in a specific tone, code in a specific internal language, or follow a strict JSON format.
Data Dynamicness	High: Data changes daily. You just update the database.	Low: Data is static. Retraining is expensive and slow.
Hallucination	Low: Grounded in retrieval.	Medium: The model can still hallucinate facts even after fine-tuning.
Cost	Low: Only inference and storage costs.	High: Requires GPUs for training and hosting custom models.

Verdict:

- Use **RAG** for factual accuracy and accessing proprietary documents.
- Use **Fine-Tuning** if you need the model to sound like a specific persona (e.g., a medical assistant) or learn a new language/code syntax.
- *Pro Tip:* Often, the best architecture is **Hybrid** (Fine-tune a model to be a good RAG reasoner, then use RAG for the facts).

5. What are the architecture patterns for customizing LLM with proprietary data?

There are three main patterns for injecting private data into an LLM application:

Pattern A: Context Injection (RAG)

- **Mechanism:** Search a database \rightarrow Stuff results into the Prompt \rightarrow Ask LLM.
- **Best For:** QA bots, Search engines, applications requiring citations.

- **Pros:** Cheapest, easiest to update, highly verifiable.

Pattern B: Small Language Model (SLM) Fine-Tuning

- **Mechanism:** Take a small open-source model (e.g., Llama-3-8B) and train it on your specific domain data (e.g., medical records).
- **Best For:** Privacy-heavy on-premise deployments, or when specific jargon/vocabulary is critical.
- **Pros:** Data never leaves your server, highly specialized performance.

Pattern C: Hybrid (RAG + Fine-Tuning)

- **Mechanism:** Fine-tune a model specifically to understand your domain's *structure* (e.g., how to read your complex financial tables), and then use RAG to feed it the actual *numbers*.
- **Best For:** Complex enterprise use cases where off-the-shelf models fail to understand the nuance of the documents.